Saiteja Utpala

saitejautpala@gmail.com saitejautpala.github.io github.com/SaitejaUtpala Google Scholar

Currently I am a researcher at MBZUAI, working on diffusion language models. Previously, I was a Research Engineer at Princeton University working on the reliability of AI agents, and a Research SDE at Microsoft Research, where I worked on the GitHub Copilot team. Before that, I worked on geometric machine learning and statistics, differential privacy, and nonparametric statistics. I have also independently contributed to several frontier evaluations, including METR, MMTEB, Humanity's Last Exam, and FrontierMath. I hold a Masters in Computer Science and Engineering from IIT Kanpur, India (2020).

Experience

MBZUAI *Researcher, Diffusion Language Models* — Abu Dhabi, UAE · Jan 2026 – Present

Princeton University *Research Engineer, AI Agent Reliability* — Princeton, USA · Sep 2025 – Dec 2025

Entelligence AI *Research SDE* — Remote, India · Mar 2025 – Sep 2025

Microsoft Research *Research SDE, GitHub Copilot Workspace* — Bangalore, India · Feb 2024 – Feb 2025

Wadhwani AI *Associate ML Scientist* — Delhi, India · Apr 2023 – Jan 2024

Geomstats *Maintainer and Lead Engineer* — Open Source · Dec 2021 – Dec 2022

UC Santa Barbara *Staff Research Associate* — California, USA · Dec 2021 – Jun 2022

Mastercard *Software Development Engineer* — Pune, India · Aug 2020 – Dec 2021

Selected Awards and Honors

2025 **UK AISI** evals bounty, £20,025, for dangerous-capability and scheming evaluations of LLM agents.

2025 **Epoch AI FrontierMath**, $1,700, for research-level math questions (invite-only).

2024 **METR** Task Bounty, $70,214, for ML R&D tasks used in METR's time-horizon evaluations.

2024 **Humanity's Last Exam**, $5,000, for a Top 50 hardest question; ranked #29 of 1,015 contributors.

2021 **ICLR 2021** Computational Geometry & Topology Challenge, $750, 2nd place.

2019 **Academic Excellence Award**, IIT Kanpur, for exceptional performance in the 2018–19 academic session.

Publications

7. **Looped Uniform Diffusion Language Models**
Saiteja Utpala, Chieh-Hsin Lai, Yuki Mitsufuji, Praneeth Vepakomma. *Preprint, 2026*.
Diffusion language models · Recurrent depth · FLOP-efficient training

6. **Towards a Science of AI Agent Reliability**
Stephan Rabanser, Sayash Kapoor, Peter Kirgis, Kangheng Liu, **Saiteja Utpala**, Arvind Narayanan. ***ICML 2026***.
AI agents · Reliability

5. **Language Agnostic Code Embeddings**
Saiteja Utpala, Alex Gu, Pin-Yu Chen. ***NAACL 2024***.
Code LLMs · Retrieval · Zero-shot alignment

4. **Locally Differentially Private Document Generation Using Zero Shot Prompting**
Saiteja Utpala, Sara Hooker, Pin-Yu Chen. ***EMNLP Findings 2023***.
LLMs · De-anonymization attacks · Private documents

3. **Shrinkage Estimation of Higher Order Bochner Integrals**
Saiteja Utpala, Bharat K. Sriperumbudur. ***Bernoulli 2023***.
Learning theory · Concentration inequalities · Functional analysis

2. **Improved Differentially Private Riemannian Optimization: Fast Sampling and Variance Reduction**
Saiteja Utpala, Andi Han, Pratik Jawanpuria, Bamdev Mishra. ***TMLR 2023***.
Differential privacy · Non-convex optimization · Riemannian optimization

1. **Differentially Private Fréchet Mean on the Manifold of SPD Matrices**
Saiteja Utpala, Praneeth Vepakomma, Nina Miolane. ***TMLR 2023***.
Differential privacy · Geometric machine learning